

Capstone Project

Customer Segmentation

Name : Dhivakar.R

Course: AI and ML(Aug2020)

Problem Statement: Grouping customers into sections based on their common characteristics is called Customer Segmentation. These clusters allow the companies to target the customers with the correct marketing message and tailor their offers for a specific group. This not only helps them boost their sales, but also helps them build customer relations and understand them in a better way.

In this project, our aim will be to perform customer segmentation on Online Retail Dataset (<https://archive.ics.uci.edu/ml/datasets/Online+Retail#>) to understand the customers. Given this dataset, our task is to:

- a) Load the dataset and perform a descriptive analysis on it (Total number of entries, the column types, unique/non-null entries for each attribute, unique stock items, visualizing various attributes using bar charts/piecharts and so on).
- b) Perform data cleaning. Specifically, given the dataset, handle the entries that either have missing information or have attribute values that are not feasible such as negative quantity.
- c) Perform data pre-processing for the required attribute fields
- d) Since this database has no additional attribute information for the customer, we will use RFM model (refer: <https://clevertap.com/blog/rfmanalysis/>) for segmentation. Modify the database to include RFM model attributes.
- e) Now once you have your database ready, perform data clustering on this dataset by assuming a fixed number of clusters.
- f) Find the optimal number of clusters that the customers can be divided into.

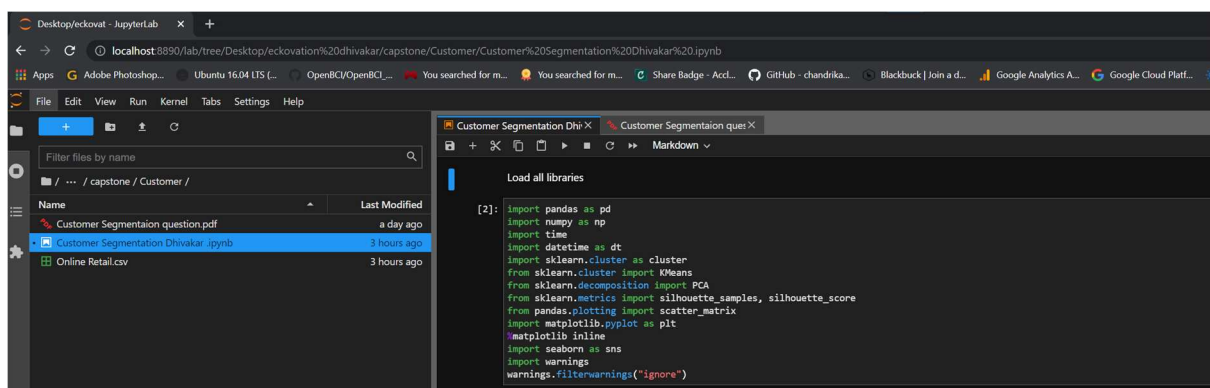
Prerequisites

What things you need to install the software and how to install them:

Python 3.6 This setup requires that your machine has latest version of python. The following url <https://www.python.org/downloads/> can be referred to download python. Once you have python downloaded and installed, you will need to setup PATH variables (if you want to run python program directly, detail instructions are below in how to run software section). To do that check this: <https://www.pythoncentral.io/add-python-to-path-python-is-not-recognized-as-an-internal-or-externalcommand/>. Setting up PATH variable is optional as you can also run program without it and more instruction are given below on this topic.

Second and easier option is to download anaconda and use its anaconda prompt to run the commands. To install anaconda check this url <https://www.anaconda.com/download/> You will also need to download and install below 3 packages after you install either python or anaconda from the steps above Sklearn (scikit-learn) numpy scipy if you have chosen to install python 3.6 then run below commands in command prompt/terminal to install these packages `pip install -U scikit-learn` `pip install numpy` `pip install scipy` if you have chosen to install anaconda then run below commands in anaconda prompt to install these packages `conda install -c scikit-learn` `conda install -c anaconda numpy` `conda install -c anaconda scipy`

Importing the libraries and loading dataset :



The screenshot shows the JupyterLab interface. On the left, the file explorer displays the directory structure: `/capstone/ Customer/`. It lists three files: `Customer Segmentation question.pdf` (modified a day ago), `Customer Segmentation Dhiavakr .ipynb` (modified 3 hours ago), and `Online Retail.csv` (modified 3 hours ago). The `Customer Segmentation Dhiavakr .ipynb` file is selected. The main code editor shows the following code:

```
[2]: import pandas as pd
import numpy as np
import time
import datetime as dt
import sklearn.cluster as cluster
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_samples, silhouette_score
from pandas.plotting import scatter_matrix
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```



The screenshot shows the JupyterLab code editor with the following code:

```
[3]: retail = pd.read_csv("Online retail.csv", encoding='cp874')
retail.head()
```

The output of the code is a preview of the first five rows of the `retail` dataset. The data is presented in a table with the following columns: `InvoiceNo`, `StockCode`, `Description`, `Quantity`, `InvoiceDate`, `UnitPrice`, `CustomerID`, and `Country`.

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	12/1/2010 8:26	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/2010 8:26	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/2010 8:26	3.39	17850.0	United Kingdom

```
[4]: retail.shape
```

```
[4]: (541909, 8)
```

```
[5]: retail.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   InvoiceNo        541909 non-null object
1   StockCode        541909 non-null object
2   Description      540455 non-null object
3   Quantity         541909 non-null int64
4   InvoiceDate       541909 non-null object
5   UnitPrice        541909 non-null float64
6   CustomerID       406829 non-null float64
7   Country          541909 non-null object
dtypes: float64(2), int64(1), object(5)
memory usage: 33.1+ MB
```

Checking the null values in the dataset

Checking the null values :

```
[7]: retail.isna().sum().sort_values(ascending=False)
```

```
[7]: CustomerID    135080
Description    1454
Country        0
UnitPrice      0
InvoiceDate    0
Quantity        0
StockCode      0
InvoiceNo      0
dtype: int64
```

```
[8]: pd.DataFrame(data = (retail.isna().sum() / retail.shape[0]) * 100, index = retail.columns, columns = ['% Null Values'])
```

```
[8]:
```

	% Null Values
InvoiceNo	0.000000
StockCode	0.000000
Description	0.268311
Quantity	0.000000
InvoiceDate	0.000000
UnitPrice	0.000000
CustomerID	24.926694
Country	0.000000

Dropping the rows with null values in customer id column

```
[9]: retail.duplicated().sum()

[9]: 5268

[11]: retail.drop_duplicates(inplace=True)
      retail.shape

[11]: (536641, 8)
```

Removing the cancelled orders from the dataset

```
Removing the cancelled orders from the dataset

[12]: retail = retail[retail['Quantity'] > 0]
      retail.shape

[12]: (526054, 8)

[13]: pd.DataFrame(data=[retail['InvoiceNo'].nunique(),retail['StockCode'].nunique(),retail['CustomerID'].nunique()],columns=['Count'],
                    index=['Number of Transactions','Number of Unique Products Bought','Number of Unique Customers'])

[13]:
```

	Count
Number of Transactions	20728
Number of Unique Products Bought	3941
Number of Unique Customers	4339

RFM Analysis

```
RFM Analysis

[16]: retail['InvoiceDate'] = retail['InvoiceDate'].astype('datetime64')
      retail['InvoiceDate'].max()

[16]: Timestamp('2011-12-09 12:50:00')

[17]: now = dt.date(2011,12,9)
      print(now)

2011-12-09

[18]: retail['Date'] = retail['InvoiceDate'].apply(lambda x: x.date())

[19]: retail.head()

[19]:
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	Date
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	2010-12-01 08:26:00	2.55	17850.0	United Kingdom	2010-12-01
1	536365	71053	WHITE METAL LANTERN	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	2010-12-01 08:26:00	2.75	17850.0	United Kingdom	2010-12-01
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01

```


[20]: recency_df = retail.groupby(by='CustomerID', as_index=False)['Date'].max()
      recency_df.columns = ['CustomerID', 'LastPurshaceDate']
      recency_df.head()

[20]:
```

	CustomerID	LastPurshaceDate
0	12346.0	2011-01-18
1	12347.0	2011-12-07
2	12348.0	2011-09-25
3	12349.0	2011-11-21
4	12350.0	2011-02-02

RFM Table:

```
Create RFM Table

[28]: rfm_df = recency_df.merge(frequency_df,on='CustomerID').merge(monetary_df,on='CustomerID')
      rfm_df.set_index('CustomerID',inplace=True)
      rfm_df.head()
```

```
[28]:
```

	Recency	Frequency	Monetary
CustomerID			
12346.0	325	1	77183.60
12347.0	2	7	4310.00
12348.0	75	4	1797.24
12349.0	18	1	1757.55
12350.0	310	1	334.40

Customer segments with RFM Model

```
Customer segments with RFM Model

[29]: pareto_cutoff = rfm_df['Monetary'].sum() * 0.8
      print("The 80% of total revenue is: ",round(pareto_cutoff,2))

The 80% of total revenue is: 7109767.12

[30]: customers_ranked = rfm_df
      customers_ranked['Rank'] = customers_ranked['Monetary'].rank(ascending=False)
      customers_ranked.head()
```

```
[30]:
```

	Recency	Frequency	Monetary	Rank
CustomerID				
12346.0	325	1	77183.60	10.0
12347.0	2	7	4310.00	335.0
12348.0	75	4	1797.24	1004.0
12349.0	18	1	1757.55	1027.0
12350.0	310	1	334.40	3097.0

```
[31]: customers_ranked.sort_values(by='Rank',ascending=True,inplace=True)
      customers_ranked.head()
```

```
[31]:
```

	Recency	Frequency	Monetary	Rank
CustomerID				
14646.0	1	74	280206.02	1.0
18102.0	0	60	259657.30	2.0
17450.0	8	46	194390.79	3.0
16446.0	0	2	168472.50	4.0
14911.0	1	201	143711.17	5.0

```
[32]: # Get top 20% of the customers
top_20_cutoff = 4339 * 20 / 100
top_20_cutoff
```

```
[32]: 867.8
```

```
[33]: # Sum the monetary values over the customer with rank <= 868
revenueByTop20 = customers_ranked[customers_ranked['Rank'] <= 868]['Monetary'].sum()
revenueByTop20
```

```
[33]: 6637300.820999999
```

Applying RFM Score Formula

```
[34]: quantiles = rfm_df.quantile(q=[0.25,0.5,0.75])
quantiles
```

```
[34]:
```

	Recency	Frequency	Monetary	Rank
0.25	17.0	1.0	306.455	1085.5
0.50	50.0	2.0	668.560	2170.0
0.75	141.5	5.0	1660.315	3254.5

Creation of RFM Segmentation Table:

Creation of RFM segmentation table

```
[36]: # Arguments (x = value, p = recency, monetary_value, frequency, d = quantiles dict)
def RScore(x,p,d):
    if x <= d[p][0.25]:
        return 4
    elif x <= d[p][0.50]:
        return 3
    elif x <= d[p][0.75]:
        return 2
    else:
        return 1
```

```
[37]: # Arguments (x = value, p = recency, monetary_value, frequency, k = quantiles dict)
def FMScore(x,p,d):
    if x <= d[p][0.25]:
        return 1
    elif x <= d[p][0.50]:
        return 2
    elif x <= d[p][0.75]:
        return 3
    else:
        return 4
```

```
[38]: # Create rfm segmentation table
rfm_segmentation = rfm_df
rfm_segmentation['R_Quartile'] = rfm_segmentation['Recency'].apply(RScore, args=('Recency',quantiles,))
rfm_segmentation['F_Quartile'] = rfm_segmentation['Frequency'].apply(FMScore, args=('Frequency',quantiles,))
rfm_segmentation['M_Quartile'] = rfm_segmentation['Monetary'].apply(FMScore, args=('Monetary',quantiles,))
```

```
[39]: rfm_segmentation.head()
```

```
[39]:
```

	Recency	Frequency	Monetary	Rank	R_Quartile	F_Quartile	M_Quartile
CustomerID							
14646.0	1	74	280206.02	1.0	4	4	4
18102.0	0	60	259657.30	2.0	4	4	4
17450.0	8	46	194390.79	3.0	4	4	4
16446.0	0	2	168472.50	4.0	4	2	4
14911.0	1	201	143711.17	5.0	4	4	4

How many customers do we have in each segment?

```
[43]: print("Best Customers: ",len(rfm_segmentation[rfm_segmentation['RFMScore']==4))
      print('Loyal Customers: ',len(rfm_segmentation[rfm_segmentation['F_Quartile']==4]))
      print("Big Spenders: ",len(rfm_segmentation[rfm_segmentation['M_Quartile']==4]))
      print('Customers at risk of churning: ', len(rfm_segmentation[rfm_segmentation['RFMScore']==3]))
      print('Almost Churned Customers: ',len(rfm_segmentation[rfm_segmentation['RFMScore']==2]))
      print('Churned Customers: ',len(rfm_segmentation[rfm_segmentation['RFMScore']==1]))
```

```
Best Customers: 455
Loyal Customers: 872
Big Spenders: 1085
Customers at risk of churning: 70
Almost Churned Customers: 10
Churned Customers: 441
```